

```
In [1]: from tensorflow.keras.datasets import mnist
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense
        from tensorflow.keras.optimizers import Adam
```

```
In [2]: # Load MNIST dataset
        (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [3]: # Preprocess data (normalize pixel values to range [0, 1])
        x_train = x_train.astype('float32') / 255
        x_test = x_test.astype('float32') / 255
```

```
In [4]: # Reshape data for input layer (28x28 pixels)
        x_train = x_train.reshape(x_train.shape[0], 28 * 28)
        x_test = x_test.reshape(x_test.shape[0], 28 * 28)
```

```
In [5]: # Convert class labels to one-hot encoded vectors
        from tensorflow.keras.utils import to_categorical
        y_train = to_categorical(y_train)
        y_test = to_categorical(y_test)
```

```
In [6]: # Define the Pure ANN model with less than 10,000 trainable parameters
        model = Sequential()
        model.add(Dense(512, activation='relu', input_shape=(28 * 28,))) # First layer with 512 neurons
        model.add(Dense(128, activation='relu')) # Second layer with 128 neurons
        model.add(Dense(10, activation='softmax')) # Output layer with 10 neurons (one for each digit class)
```

```
In [7]: # Compile the model
        model.compile(loss='categorical_crossentropy', optimizer=Adam(learning_rate=0.001), metrics=['accuracy'])
```

```
In [8]: # Train the model
model.fit(x_train, y_train, epochs=10, batch_size=128, validation_data=(x_test, y_test))
```

```
Epoch 1/10
469/469 [=====] - 4s 6ms/step - loss: 0.2393 - accuracy: 0.9304 - val_loss: 0.1190 - val_accuracy: 0.9643
Epoch 2/10
469/469 [=====] - 3s 6ms/step - loss: 0.0895 - accuracy: 0.9726 - val_loss: 0.0805 - val_accuracy: 0.9744
Epoch 3/10
469/469 [=====] - 3s 6ms/step - loss: 0.0550 - accuracy: 0.9829 - val_loss: 0.0694 - val_accuracy: 0.9786
Epoch 4/10
469/469 [=====] - 3s 6ms/step - loss: 0.0391 - accuracy: 0.9880 - val_loss: 0.0666 - val_accuracy: 0.9794
Epoch 5/10
469/469 [=====] - 3s 6ms/step - loss: 0.0278 - accuracy: 0.9909 - val_loss: 0.0712 - val_accuracy: 0.9795
Epoch 6/10
469/469 [=====] - 3s 6ms/step - loss: 0.0218 - accuracy: 0.9935 - val_loss: 0.0672 - val_accuracy: 0.9808
Epoch 7/10
469/469 [=====] - 3s 6ms/step - loss: 0.0191 - accuracy: 0.9936 - val_loss: 0.0716 - val_accuracy: 0.9813
Epoch 8/10
469/469 [=====] - 3s 7ms/step - loss: 0.0140 - accuracy: 0.9956 - val_loss: 0.0717 - val_accuracy: 0.9811
Epoch 9/10
469/469 [=====] - 3s 6ms/step - loss: 0.0146 - accuracy: 0.9953 - val_loss: 0.1071 - val_accuracy: 0.9749
Epoch 10/10
469/469 [=====] - 3s 6ms/step - loss: 0.0128 - accuracy: 0.9958 - val_loss: 0.0878 - val_accuracy: 0.9792
```

```
Out[8]: <keras.src.callbacks.History at 0x1b63652f700>
```

```
In [9]: # Evaluate the model on test data
test_loss, test_acc = model.evaluate(x_test, y_test)
print('Test accuracy:', test_acc)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.0878 - accuracy: 0.9792
Test accuracy: 0.979200005531311
```

```
In [10]: # Count the total number of trainable parameters
total_params = model.count_params()
print('Total trainable parameters:', total_params)
```

```
Total trainable parameters: 468874
```